

Forecast Audit

Day 1A: Capabilities, Diffusion, and Forecasting

Purpose. Evaluate a dated AI scenario without reducing the exercise to a binary verdict. Distinguish observed facts, explicit predictions, hidden assumptions, and claims that remain untestable as of mid-2026.

Readings

Mustafa Suleyman, *The Coming Wave*, Ch. 1; Stanford HAI, “12 Takeaways from the 2026 AI Index”; selected sections of *AI 2027*. The optional Rothman essay and the authors’ later *AI 2040* scenario may be used as counterpoints.

Team Assignments

Each team audits one part of the scenario:

- Team 1: agent reliability and “Stumbling Agents”
- Team 2: coding automation and AI-assisted research
- Team 3: scaling, compute, and capability progress
- Team 4: organizational adoption and economic integration
- Team 5: diffusion, containment, and competitive pressure
- Team 6: geopolitics and international response

Audit Rules

1. Use an evidence cutoff of **July 2026**. Do not score later claims.
2. Break broad claims into propositions that could in principle be observed.
3. Classify each proposition as **consistent**, **in tension**, or **unresolved**.
4. A claim is not validated merely because a related event occurred.
5. Identify the causal assumptions connecting the observed event to the forecast.

Team Worksheet

Complete at least four rows.

Scenario claim	Evidence through mid-2026	Status	Assumptions or missing evidence

Required Conclusion

Prepare a three-minute briefing that identifies:

- one near-term prediction that is reasonably consistent with developments;
- one prediction or implied trajectory that is in tension with developments;
- one important proposition that cannot yet be evaluated;
- the single assumption doing the most work in your section of the scenario.

Closing Question

When authors revise the expected timing of automated AI research, is that evidence of responsible updating, weak forecasting foundations, or both?