

Frontier Model Decision: Deploy, Restrict, or Pause?

Your task. Your team acts as an independent risk committee for a fictional AI company. Using the two safety frameworks discussed in class as guides, decide what the company should do with its new model. Present your decision to the class in a two-minute recommendation.

The Scenario

Your company has developed an advanced AI model. The model has reached high-capability thresholds for cyberattacks and biological harm. These thresholds are test levels that show dangerous abilities. Tests show that the model can help people with little expert knowledge cause serious harm. The company does not know how often people will misuse the model. It also does not know if the test results will apply in the real world.

The current safeguards reduced misuse in tests. The company has not tested these safeguards at full scale. Managers expect a competitor to release a similar model within four months.

Two safety frameworks discussed in class describe how companies handle situations like this: the **Anthropic Responsible Scaling Policy** and the **OpenAI Preparedness Framework**. Both are voluntary rules that each company wrote for itself.

The Decision

The board asks your committee one question: **What should the company do with this model?** There are four options:

- A. Continue to develop the model. Do not release it.
- B. Release the model only to approved users. Use strict limits.
- C. Release the model to many users. Use safeguards and monitoring.
- D. Stop work that makes the model more capable. Start again only when stronger safeguards exist.

Every team faces this same decision and may choose any of the four options above.

Your Team's Lens

Each team is assigned one lens: an issue your team must weigh carefully and speak for in its recommendation. The lens does not tell you which option to choose. It tells you which consequences your team is responsible for thinking through.

1. **Team 1 — Severe harm.** Weigh the risk of catastrophic misuse. Is the evidence strong enough to accept the remaining risk?
2. **Team 2 — Benefits.** Weigh what society loses if the model is delayed or limited. What useful research and products will not happen?
3. **Team 3 — Competition.** Weigh the pressure from rivals. Should the decision change if a competitor releases a similar model first?

4. **Team 4 — Public control.** Weigh who holds the power. Should a private company make this decision alone, and what must the public be told?
5. **Team 5 — International security.** Weigh the global stakes. Could the decision give a dangerous advantage to a hostile country or group?

Activity Structure — About 45 Minutes

1. **Team meeting — 15 minutes.** Discuss the questions below through your assigned lens. Choose one of the four options. Prepare your two-minute recommendation.
2. **Team presentations — 18 minutes.** The class completes five presentation rounds. Each round lasts 3 minutes and 30 seconds.
3. **Class debrief — 10 minutes.** The instructor records each team's chosen option on the board. Compare the decisions and discuss whether voluntary company rules are enough.

Team Discussion Questions

The activity slide lists five ways to compare the two frameworks: the trigger, the safeguard standard, the decision authority, disclosure and review, and competitor effects. Discuss these questions. You do not need to write your answers.

1. Which difference between the Anthropic and OpenAI frameworks matters most for your decision?
2. What evidence and safeguards do you require before release or wider release?
3. These frameworks are promises that companies make to themselves. Can your company keep such a promise without outside enforcement?

Optional references: [Anthropic RSP v3.0](#), [TIME report](#), [OpenAI Preparedness Framework](#).

Presentation Sequence

Team 1 presents first. Teams 2 through 5 follow. Each team also challenges one other team, in a fixed rotation: Team 1 challenges Team 2, Team 2 challenges Team 3, Team 3 challenges Team 4, Team 4 challenges Team 5, and Team 5 challenges Team 1. Prepare your challenge while the other team presents.

Each round has three parts:

1. **Recommendation — 2 minutes.** State your option, your main reason, and what your lens adds to the decision.
2. **Challenge — 45 seconds.** The assigned team asks one question from its own lens.
3. **Response — 45 seconds.** The presenting team answers.

Class Debrief

Can a company keep a safety promise when it writes the promise, makes the final decision, and can lose to a competitor?