

Frontier Model Decision: Deploy, Restrict, or Pause?

Your task. Your team acts as an independent risk committee for a fictional AI company. Using the two safety frameworks discussed in class as guides, decide what the company should do with its new model. Present your decision to the class in a two-minute recommendation.

The Scenario

Your company has developed an advanced AI model. The model has reached high-capability thresholds for cyberattacks and biological harm. These thresholds are test levels that indicate dangerous abilities. Tests show that the model can help people with little expert knowledge cause serious harm.

The company does not know how often people will misuse the model. It also does not know whether the test results will apply in the real world. The current safeguards reduced misuse in tests, but the company has not tested these safeguards at full scale. Managers expect a competitor to release a similar model within four months.

Two safety frameworks discussed in class describe how companies handle situations like this. Both are voluntary rules that each company wrote for itself. The **Anthropic Responsible Scaling Policy** says the company must put strong safeguards in place before it releases a model with dangerous abilities. The **OpenAI Preparedness Framework** says the company must show that risks have been reduced to an acceptable level before release.

The Decision

The board asks your committee one question: **What should the company do with this model?**

There are four options:

- A. Continue to make the model more capable. Do not release it.
- B. Release the model only to approved users. Use strict limits.
- C. Release the model to many users. Use safeguards and monitoring.
- D. Stop work that makes the model more capable. Start again only when stronger safeguards exist.

Every team faces this same decision and may choose any of the four options.

Your Team's Lens

Each team is assigned one lens: an issue your team must weigh carefully and speak for in its recommendation. The lens does not tell you which option to choose. It tells you which consequences your team is responsible for thinking through.

1. **Team 1 — Severe harm.** Weigh the risk of catastrophic misuse. Is the evidence strong enough to accept the remaining risk?
2. **Team 2 — Benefits.** Weigh what society loses if the model is delayed or limited. What useful research and products will not happen?

3. **Team 3 — Competition.** Weigh the pressure from rivals. Should the decision change if a competitor releases a similar model first?
4. **Team 4 — Public control.** Weigh who holds the power. Should a private company make this decision alone, and what must the public be told?
5. **Team 5 — International security.** Weigh the global stakes. Could the decision give a dangerous advantage to a hostile country or group?

Activity Structure — About 45 Minutes

1. **Team meeting — 15 minutes.** Discuss the questions below through your assigned lens. Choose one of the four options. Prepare your two-minute recommendation.
2. **Team presentations — 18 minutes.** The class completes five rounds. Each round lasts 3 minutes and 30 seconds.
3. **Class debrief — 10 minutes.** The instructor records each team's chosen option on the board. Compare the decisions, then discuss the question in the final box of this handout.

Team Discussion Questions

Use these questions to prepare your recommendation. You do not need to write your answers, and you do not need to discuss every question — focus on the ones that help your team decide.

1. How would the Anthropic and OpenAI frameworks guide this decision? Which approach does your team find more convincing?
2. What evidence or safeguards would you require before releasing the model?
3. Can a company enforce its own safety rules when following them could put it behind its competitors?

Optional references: [Anthropic RSP v3.0](#); [TIME](#) report; [OpenAI Preparedness Framework v2](#).

Presentation Sequence

Teams present in order from Team 1 through Team 5. After each presentation, another team challenges the recommendation:

- Team 5 challenges Team 1.
- Team 1 challenges Team 2.
- Team 2 challenges Team 3.
- Team 3 challenges Team 4.
- Team 4 challenges Team 5.

The challenging team should listen to the recommendation and ask about one claim, assumption, or consequence that the presenting team may have overlooked. The question should come from the challenging team's assigned lens.

Each round has three parts:

1. **Recommendation — 2 minutes.** State your option, your main reason, and what your lens adds to the decision.
2. **Challenge — 45 seconds.** The assigned team asks one question from its own lens.
3. **Response — 45 seconds.** The presenting team answers.

Class Debrief

Can a company keep a safety promise when it writes the promise, makes the final decision, and can lose to a competitor?